

Inside the AI Conversation: Sub-Category Drill-Downs on TikTok, 2026

Pairie Koh

Andrew B. Hall

June 1, 2026

Abstract

This memo opens up the top-level theme bars from our AI-sentiment corpus and asks what is *inside* each one. We work in two tiers. **Tier 1** uses fields already in the corpus—the editorial *register* of each video and the seed search keyword that surfaced it—to characterise each theme without any new classification. **Tier 2** runs a second, independent pass with Claude (Haiku 4.5, via the Message Batches API) that assigns every classified video to one theme-specific sub-category. Everything here is **TikTok only** and restricted to videos posted in **2026**, matching the main bar charts. The recurring lesson is that the public’s AI conversation is concrete and material in ways the elite conversation is not: the deepfake fear is about *scams*, not elections; the jobs panic is generalised dread and layoff news, not a “white-collar bloodbath”; and the loudest enthusiasm is mostly marketing, while the most genuine enthusiasm is for AI *companionship*.

Contents

1	Method and scope	1
2	Tier 1: register and keyword structure	2
2.1	How organic is each positive theme?	2
2.2	Seed-keyword breakdown per theme	2
3	Tier 2: independent sub-categories	5
3.1	Backlash	5
3.2	Positive	5
3.3	Selected themes in detail	6
4	What this adds to the argument	9

1 Method and scope

The drill-downs sit on top of the corpus described in the main write-up: AI-related TikTok videos collected via Bright Data, filtered and theme-classified with Claude. We re-use the *canonical* row filter from the published bar charts (backlash = `_backlash_final=YES` with a primary topic; positive = pollution dropped, `model_stan/general_hype` secondaries promoted, `_theme` valid), so the sub-category totals add up to exactly the same bars readers already saw.

- **Tier 1 (no new classification).** The *register* field (`genuine_sentiment`, `creator_promo`, `creative_showcase`, `news_explainer`, `brand_official`) and the seed `_search_keyword` are already attached to each row.
- **Tier 2 (new sub-classification).** A second Claude pass assigns each video to one of a short, theme-specific menu of sub-categories. This is an *independent* judgement of content—not the query that found the video.

Caveats. (i) This is TikTok only; the YouTube halves of the corpus are not included here. (ii) The keyword breakdowns (Tier 1) reflect the seed query, not content, and are skewed by how many seeds we used per theme—read them as colour, not measurement. (iii) In the 2026 cut, the existential-risk ($n = 32$) and breakthrough-science ($n = 20$) sub-cells are too thin for headline percentages.

2 Tier 1: register and keyword structure

2.1 How organic is each positive theme?

Splitting each positive theme by editorial register is the single most consequential Tier 1 result, because it complicates any “AI is well loved” reading of the raw volume. The themes that dominate by count are the *least* organic: career/productivity content is 60% creator promotion (course and tool funnels) and only 32% genuine sentiment, and synthetic-media play is overwhelmingly creative showcase with just 8% expressing an actual opinion. The genuine enthusiasm is concentrated in the *smaller* themes—AI companionship (68% genuine) and model fandom (85% genuine). Breakthrough science is three-quarters reposted news, not public excitement.




Figure 1: Register split of each positive theme. Green is genuine sentiment; orange and red are promotional. The high-volume themes (top) are the most promotional; the most genuine enthusiasm sits in AI companionship and model fandom.

2.2 Seed-keyword breakdown per theme

The keyword views give a fast, if rough, sense of what sits under each bar. On the backlash side they confirm two things that matter for the narrative: within the environment theme, “AI data centers” is the single largest seed, and the art theme is dominated by movement hashtags (#noAI, #antiAI, “stop AI art”).

AI backlash — top seed keywords per theme (proxy sub-categories) · TikTok 2026




Figure 2: Top seed search keywords within each backlash theme (proxy sub-categories). Read as colour, not measurement.

AI positive — top seed keywords per theme (proxy sub-categories) · TikTok 2026




Figure 3: Top seed search keywords within each positive theme. “Vibe coding” dominates career content; “Suno” leads creative tools; the companion theme spans “my AI bestie,” “AI girlfriend,” and `character.ai`.

3 Tier 2: independent sub-categories

3.1 Backlash

The backlash drill-down delivers the memo’s sharpest mass-versus-elite contrasts. **Deepfakes** are feared as *scams*, not as a threat to democracy: 42% scam/fraud and 42% general misinformation, with political/election content at just 2%. **Jobs** anxiety is generalised—36% “all-jobs panic” plus 27% reaction to layoff news and 17% service/blue-collar—rather than the specifically white-collar story (~5%) that dominates elite commentary. The **art** backlash is not scattered grievance but an organised movement: 66% of it is #noAI activism and identity content. Within **environment**, water usage (30%) edges out data centers (21%) as the most-named resource.




Figure 4: Sub-categories within each backlash theme, TikTok 2026. Note the near-absence of election framing inside deepfakes and the small share of named white-collar inside jobs.

3.2 Positive

On the positive side, the “career win” is mostly a *tooling demo*: vibe coding (28%) and ChatGPT-workflow tips (27%) dwarf actually landing a job (4%). AI companionship is broad rather than narrowly romantic, spreading evenly across romantic, friendship, emotional-support, and roleplay use. Model fandom is Claude-led (51% of stan content).




Figure 5: Sub-categories within each positive theme, TikTok 2026.

3.3 Selected themes in detail




Figure 6: Jobs/displacement. Generalised panic and layoff news dominate; named white-collar displacement is a small slice.




Figure 7: Art/creative theft. The theme is overwhelmingly organised #noAI activism rather than individual job-loss complaints.




Figure 8: Environment/energy. Water usage leads, narrowly ahead of data centers and electricity/grid.




Figure 9: Career/productivity. The positive “career win” is mostly tool demonstrations, not reported employment outcomes.




Figure 10: AI companion. Affective use is broad—romantic, friendship, emotional support, and roleplay each take roughly a quarter.

4 What this adds to the argument

1. **The deepfake fear is about being scammed, not about elections.** The elite frame (deepfakes endanger democracy) is almost absent; the public frame is fraud and slop.
2. **The jobs panic is not “white-collar.”** It is generalised dread plus reactions to layoff headlines, which cuts against the specifically white-collar story told in elite policy memos.
3. **The art backlash is organised.** Two-thirds of it is movement activism, suggesting a coordinated constituency rather than diffuse complaint.
4. **Genuine enthusiasm is for companionship, not productivity.** Once promotion and showcase content are removed, the most heartfelt pro-AI content is about companions and model loyalty.
5. **Both elite poles are nearly invisible.** Existential risk and breakthrough science remain rounding errors even at the sub-category level.